

# Introduction to Deep Learning

24 Aug 2016

Vineeth N Balasubramanian  
Indian Institute of Technology, Hyderabad




# Relevance

## AlphaGo: The Recent Sensation




# Machine Learning (vs) Artificial Intelligence

## Introduction




# Machine Learning (vs) Artificial Intelligence

## Introduction




Deep learning: A sub-area of machine learning, that is today understood as representation learning



# Deep Learning

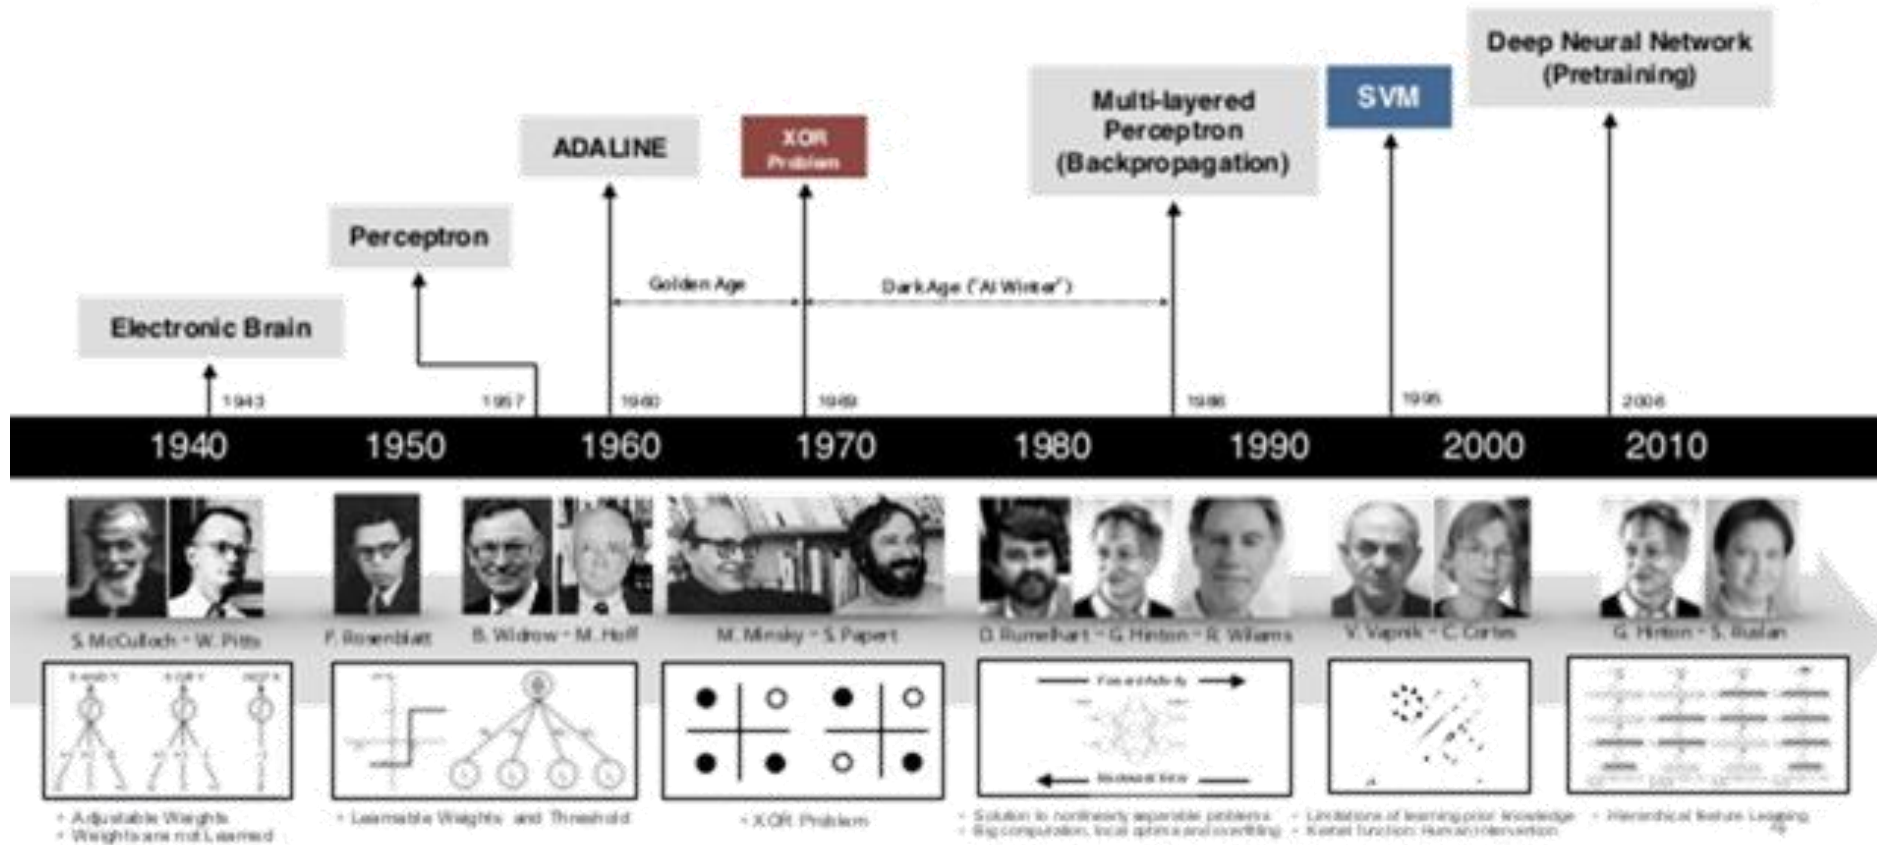
## Introduction

- Rebirth of neural networks
- Inspired by the human brain (networks of neurons)



# Deep Learning

## History






# Deep Learning

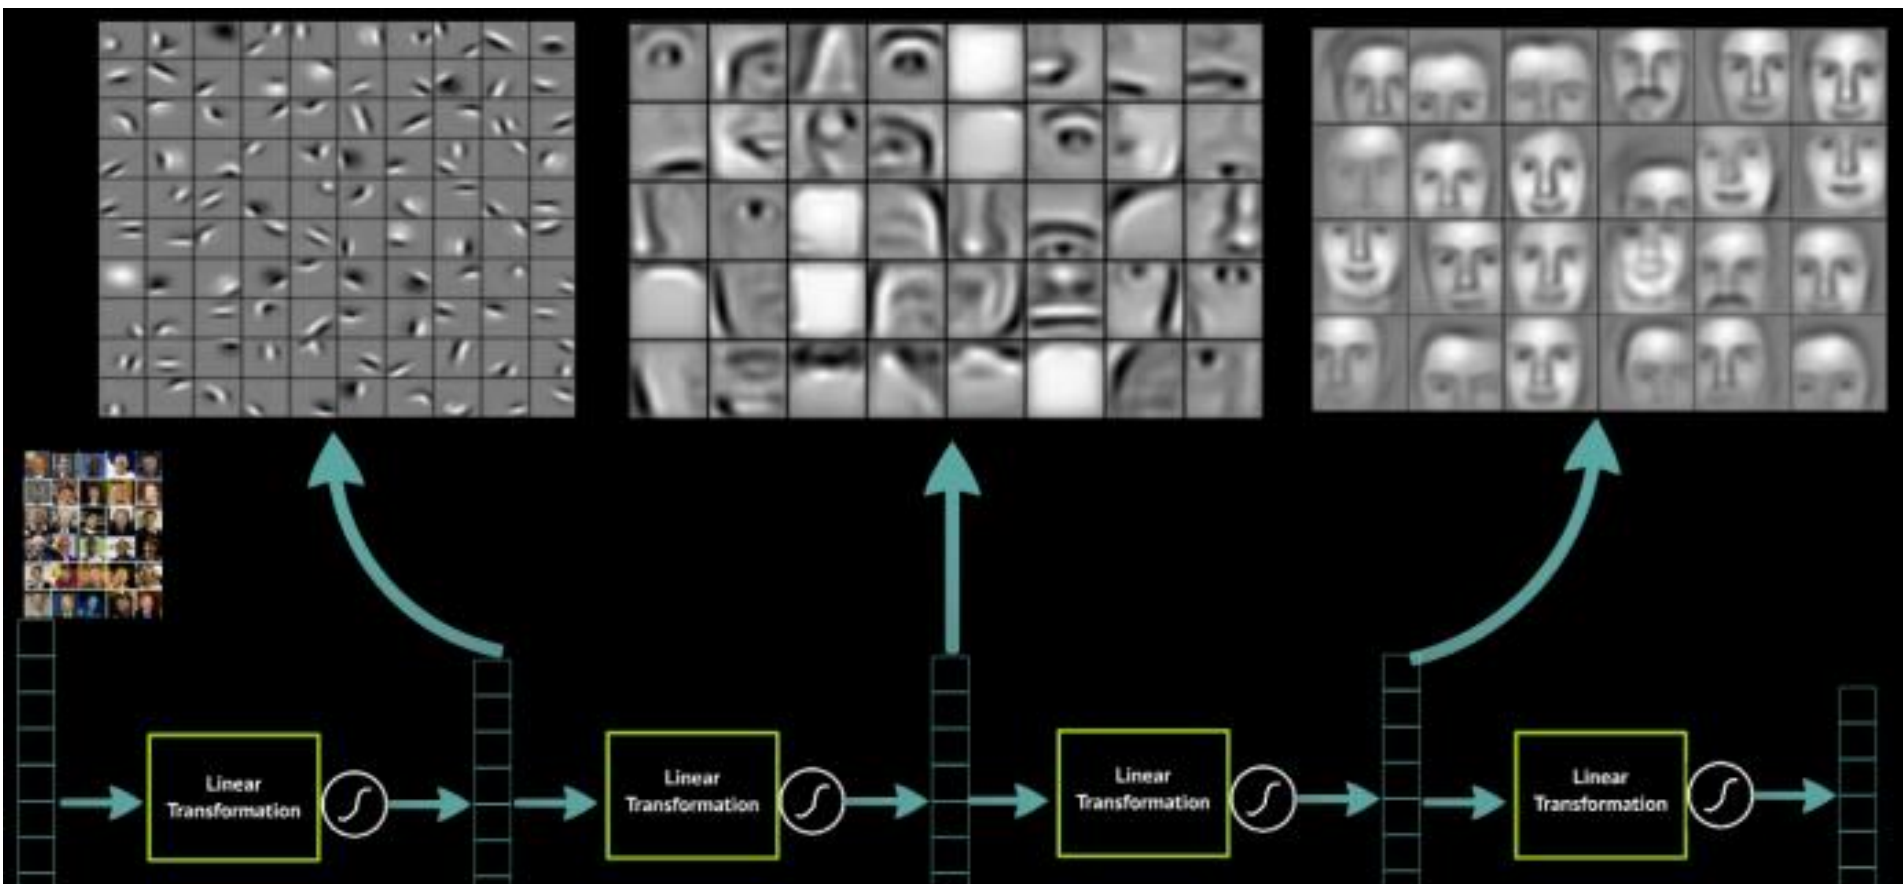
## Why is it successful?

- Learns representations of data that are useful (Other ML algorithms are “shallow”)
- Similar to the human brain
- Then, why was it not successful earlier (in the 90s)
  - Computational power
  - Data power



# Deep Learning


Why is it successful?





# Deep Learning


How do they learn?



Source: Prof Corne, Heriot-Watt University, UK

# Deep Learning

How do they learn?




Source: Prof Corne, Heriot-Watt University, UK

# Deep Learning

How do they learn?

*A dataset*

<i>Fields</i>	<i>class</i>
1.4 2.7 1.9	0
3.8 3.4 3.2	0
6.4 2.8 1.7	1
4.1 0.1 0.2	0
etc ...	



Source: Prof Corne, Heriot-Watt University, UK

# Deep Learning

How do they learn?

*Training data*

<i>Fields</i>	<i>class</i>
---------------	--------------

1.4 2.7 1.9	0
-------------	---


3.8 3.4 3.2	0
-------------	---

6.4 2.8 1.7	1
-------------	---

4.1 0.1 0.2	0
-------------	---

etc ...

Initialise with random weights



Source: Prof Corne, Heriot-Watt University, UK

# Deep Learning

How do they learn?

*Training data*

*Fields* *class*

1.4 2.7 1.9 0


3.8 3.4 3.2 0

6.4 2.8 1.7 1

4.1 0.1 0.2 0

etc ...

Present a training pattern



Source: Prof Corne, Heriot-Watt University, UK



# Deep Learning

How do they learn?

*Training data*

*Fields* *class*

1.4 2.7 1.9 0


3.8 3.4 3.2 0

6.4 2.8 1.7 1

4.1 0.1 0.2 0

etc ...

Feed it through to get output



Source: Prof Corne, Heriot-Watt University, UK

# Deep Learning

How do they learn?

*Training data*

*Fields* *class*

1.4 2.7 1.9 0


3.8 3.4 3.2 0

6.4 2.8 1.7 1

4.1 0.1 0.2 0

etc ...

Compare with target output



Source: Prof Corne, Heriot-Watt University, UK

# Deep Learning

How do they learn?

*Training data*

***Fields*** ***class***

1.4 2.7 1.9 0


3.8 3.4 3.2 0

6.4 2.8 1.7 1

4.1 0.1 0.2 0

etc ...

Adjust weights based on error



Source: Prof Corne, Heriot-Watt University, UK

# Deep Learning

How do they learn?

*Training data*

***Fields***                      ***class***

1.4 2.7 1.9              0


3.8 3.4 3.2              0

6.4 2.8 1.7              1

4.1 0.1 0.2              0

etc ...

**Present a training pattern**



Source: Prof Corne, Heriot-Watt University, UK

# Deep Learning

How do they learn?

*Training data*

<i>Fields</i>	<i>class</i>
---------------	--------------

1.4	2.7	1.9	0
-----	-----	-----	---


3.8	3.4	3.2	0
-----	-----	-----	---

6.4	2.8	1.7	1
-----	-----	-----	---

4.1	0.1	0.2	0
-----	-----	-----	---

etc ...

Feed it through to get output



Source: Prof Corne, Heriot-Watt University, UK



# Deep Learning

How do they learn?

*Training data*

***Fields***                      ***class***

1.4 2.7 1.9                      0


3.8 3.4 3.2                      0

6.4 2.8 1.7                      1

4.1 0.1 0.2                      0

etc ...

**Compare with target output**



Source: Prof Corne, Heriot-Watt University, UK

# Deep Learning

How do they learn?

*Training data*

***Fields***                      ***class***

1.4 2.7 1.9                      0


3.8 3.4 3.2                      0

6.4 2.8 1.7                      1

4.1 0.1 0.2                      0

etc ...

**Adjust weights based on error**



Source: Prof Corne, Heriot-Watt University, UK

# Deep Learning

## How do they learn?

Source: Prof Corne, Heriot-Watt University, UK

*Training data*

**Fields** **class**

1.4 2.7 1.9 0

3.8 3.4 3.2 0


6.4 2.8 1.7 1

4.1 0.1 0.2 0

etc ...

Repeat this thousands, maybe millions of times – each time taking a random training instance, and making slight weight adjustments, reduce the error


And so on ....



Called “Gradient Descent”


# The decision boundary perspective...

Initial random weights



# The decision boundary perspective...


Present a training instance / adjust the weights






# The decision boundary perspective...

Present a training instance / adjust the weights




# The decision boundary perspective...

Present a training instance / adjust the weights




# The decision boundary perspective...

Present a training instance / adjust the weights

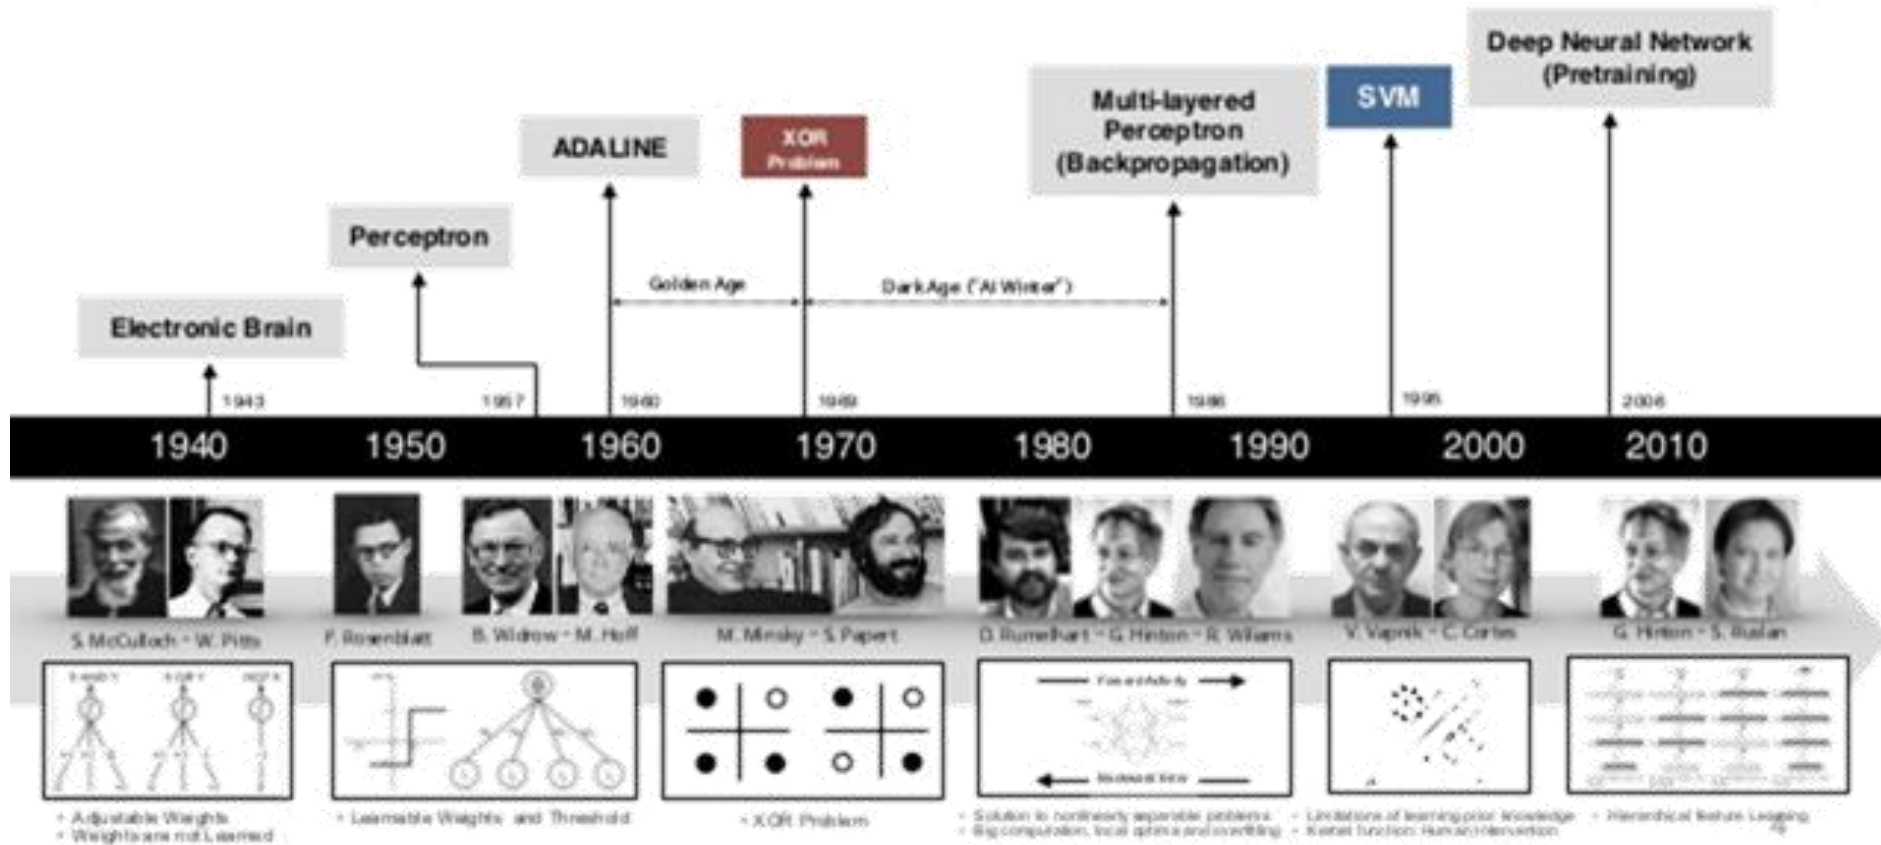


# The decision boundary perspective...

Eventually ....



# History of Neural Networks





# Neural Networks Training: Backpropagation

## Multi-Layer Perceptrons

- Extension of perceptrons to multiple layers
- 1. **Initialize** network with **random** weights
- 2. **For all** training cases (**called examples**):
  - **a.** Present training inputs to network and calculate output
  - **b.** For all layers (starting with output layer, back to input layer):
    - i. Compare **network output** with **correct output** (error function)
    - ii. **Adapt weights** in current layer

# Neural Networks Training: Backpropagation

## Multi-Layer Perceptrons

- Method for **learning weights** in feed-forward (FF) nets
- Can't use Perceptron Learning Rule
  - no **teacher values** are possible for **hidden units**
- Use **gradient descent** to minimize the error
  - **propagate deltas** to **adjust for errors backward from outputs**  
to hidden layers  
to inputs

# Neural Networks Training: Backpropagation

## Multi-Layer Perceptrons

- The idea of the algorithm can be summarized as follows :
  1. Computes the **error term for the output units** using the observed error.
  - 2. From output layer, **repeat**
    - propagating the error term back to the previous layer and updating the weights between the two layers until the earliest hidden layer is reached.

# Neural Networks Training: Backpropagation

## Multi-Layer Perceptrons

- Initialize weights (typically random!)
- Keep doing epochs
  - **For each** example **e** in training set do
    - **forward pass** to compute
      - $O = \text{neural-net-output}(\text{network}, e)$
      - $\text{miss} = (T - O)$  at each output unit
    - **backward pass** to calculate deltas to weights
    - update all weights
    - end
  - until **tuning set error stops improving**


Forward pass

Backward pass explained in next slide

# Neural Networks Training: Backpropagation

## Gradient Descent

- Think of the  $N$  weights **as a point** in an  $N$ -dimensional space
- Add a **dimension** for the observed error
- Try to **minimize your position** on the “error surface”



# Neural Networks Training: Backpropagation

Compute  
deltas

Gradient Descent

- Trying to make **error decrease the fastest**
- **Compute:**
  - $\text{Grad}_E = [dE/dw_1, dE/dw_2, \dots, dE/dw_n]$
- **Change  $i$ -th weight by**
  - $\text{delta}_{w_i} = -lr * dE/dw_i$
  - $lr$  or  $\alpha$  learning rate: generally starts with  $1E-1$  or  $10^{-1}$
- We need a **derivative!**
- Activation function must be **continuous**, differentiable, non-decreasing, and easy to compute

Derivatives of error for weights

# Neural Networks Training: Backpropagation

## Updating Hidden-to-Output

- We have **teacher supplied** desired values

- $$\text{delta}_{wji} = \alpha * a_j * (T_i - O_i) * g'(in_i)$$
$$= \alpha * a_j * (T_i - O_i) * O_i * (1 - O_i)$$

– for sigmoid the derivative is,  $g'(x) = g(x) * (1 - g(x))$

*alpha*

Here we have general formula with derivative, next we use for sigmoid

miss

derivative



# Making Choices

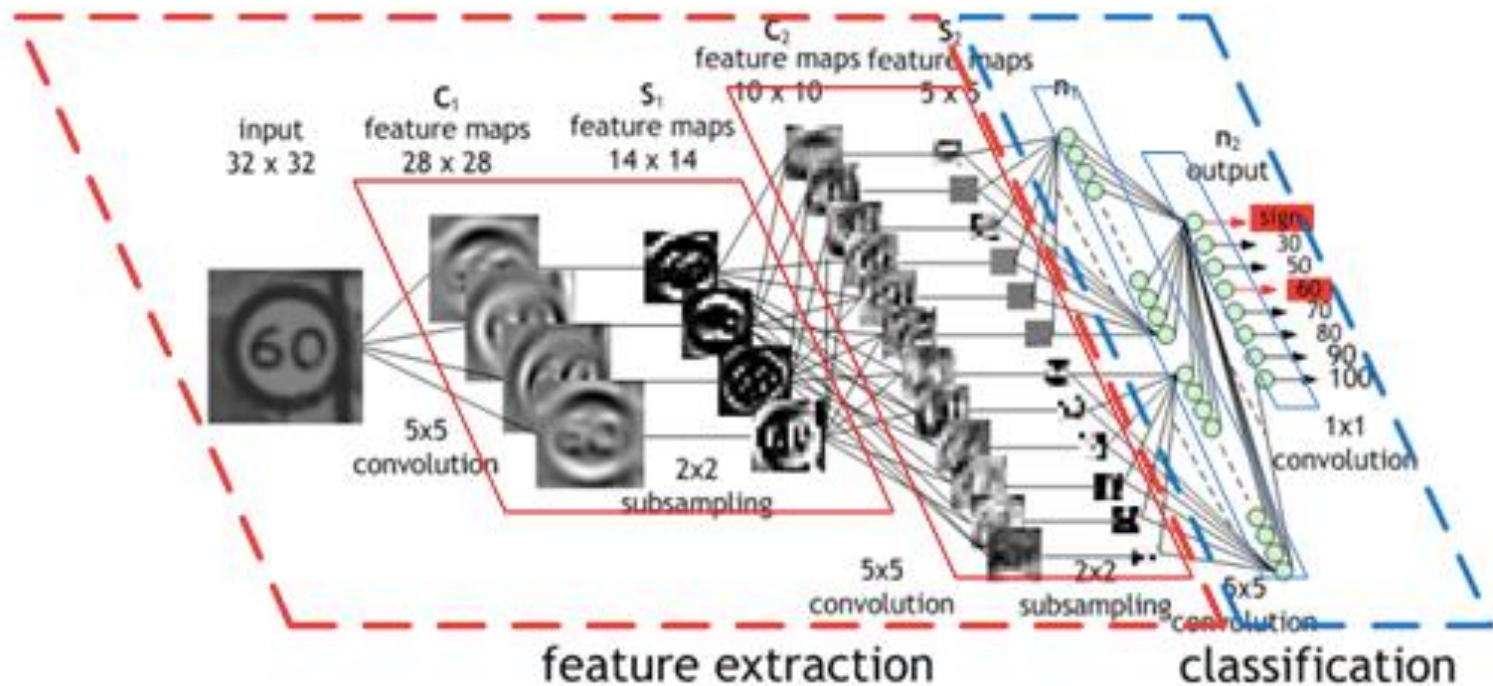
## Backpropagation

- How do we pick  **$\alpha$  or  $lr$** ?
  - **Tuning** set, or
  - **Cross validation**, or
  - **Small** for slow, conservative learning
    - for Deep learning typically is  $1e-1$
- How many hidden layers?
  - **Too few** ==> can't learn
  - Too many ==> poor generalization
- How big a training set?
  - Determine your **target error rate,  $e$** ; **Success rate** is  $1 - e$
  - Typical training set approx.  **$n/e$** , where  **$n$**  is the number of weights in the net
  - Example:
    - $e = 0.1$ ,  $n = 80$  weights; Training set **size 800**

# Deep Learning Architectures

## Variants


### Convolutional Neural Networks for Image and Video Understanding



# Deep Learning Architectures

## Variants

### Recurrent Neural Networks for Time Series and Sequence Data Understanding



# Deep Learning

## Applications and Successes

- AlexNet (Object Recognition): The network that catapulted the success of deep learning in 2012



# Stepping to even higher levels of intelligence

## From generation to imagination

- Deep dreaming/hallucination
  - <https://www.youtube.com/watch?v=oyxSerkkP4o>





# Stepping to even higher levels of intelligence

From generation to imagination



Horizon



Towers & Pagodas



Trees



Buildings



Leaves




Birds & Insects

# Stepping to even higher levels of intelligence

## From generation to imagination

- Deep colorization



# What's cutting edge?

## Recent Trends

- Limitless applications
  - Any application where you would like to learn from data
  - When do you choose deep networks against shallow ML models?
- Neural network compression
- Specialized hardware for deep learning
  - Movidius (<http://www.movidius.com/>) - deep learning accelerator



# More?

Where to look?

- One-stop shop
  - <https://github.com/ChristosChristofidis/awesome-deep-learning>
- Check this out for hours of fun and amazement
  - <http://fastml.com/deep-nets-generating-stuff/>
- Books (on Deep Learning)
  - <http://www.deeplearningbook.org>
  - <http://neuralnetworksanddeeplearning.com/>
- Programming
  - Theano/Pylearn2, Caffe, Torch (used in Google, Facebook and other companies), TensorFlow – recent initiative of Google, Keras